FraudWatch — Real-Time Fraud Detection using Big Data Analytics and Machine Learning

Draft report — 2026-09-18. This draft is written against the project's own evidence files. It is labelled a draft because the final institutional template has not been supplied; when it arrives the content below should be re-fitted to that template rather than rewritten from memory. Every number in this report is traceable to a file in the repository, and each section says plainly where a claim comes from and what was *not* measured.

---

1. Cover

Project title	FraudWatch — Real-Time Fraud Detection using Big Data Analytics and Machine Learning
Group	Group 4
Members	Hardik Hinduja · Avinash Gehi · Siddharth Dey · Sahil Deshmukh
Submission date	Monday, 21 September 2026
Deployed sites	FraudWatch Live · FraudWatch Lab (§8)
Repository	monorepo — apps, services, packages, infra, ml, docs, reports
Report status	Draft, evidence-based; see §10 for the open items

---

2. Abstract, problem statement, objectives, CA1/CA2 mapping

Abstract

FraudWatch is a working fraud-detection system, not a mockup. Transactions from the IEEE-CIS Fraud Detection dataset are replayed onto Apache Kafka, enriched with per-account streaming features by Apache Flink and Apache Spark Structured Streaming, scored by a LightGBM model behind one shared HTTP scoring service, persisted to PostgreSQL, and streamed to two React websites. The whole backend runs on a single Google Compute Engine VM that wakes when a visitor arrives and stops itself after 45 idle minutes, so the deployment stays inside Google Cloud Free Trial credit. The system has been exercised end to end on GCP: 500 records processed by both engines with bit-identical scores, a measured late event path, a real TaskManager kill with checkpoint recovery, and a full sleep/wake/sleep lifecycle.

Problem statement

Card fraud must be caught while a transaction is still in flight, which makes it a streaming problem rather than a batch problem. Two properties make it hard. First, fraud is rare: in the held-out test split used here

only 3.48% of transactions are fraudulent, so accuracy is a meaningless headline and precision/recall trade-offs dominate. Second, the arrival order is not guaranteed: events can arrive late or out of order, and a system that silently drops them reports a cleaner result than it earned. FraudWatch is built to make both of these visible rather than to hide them.

Objectives

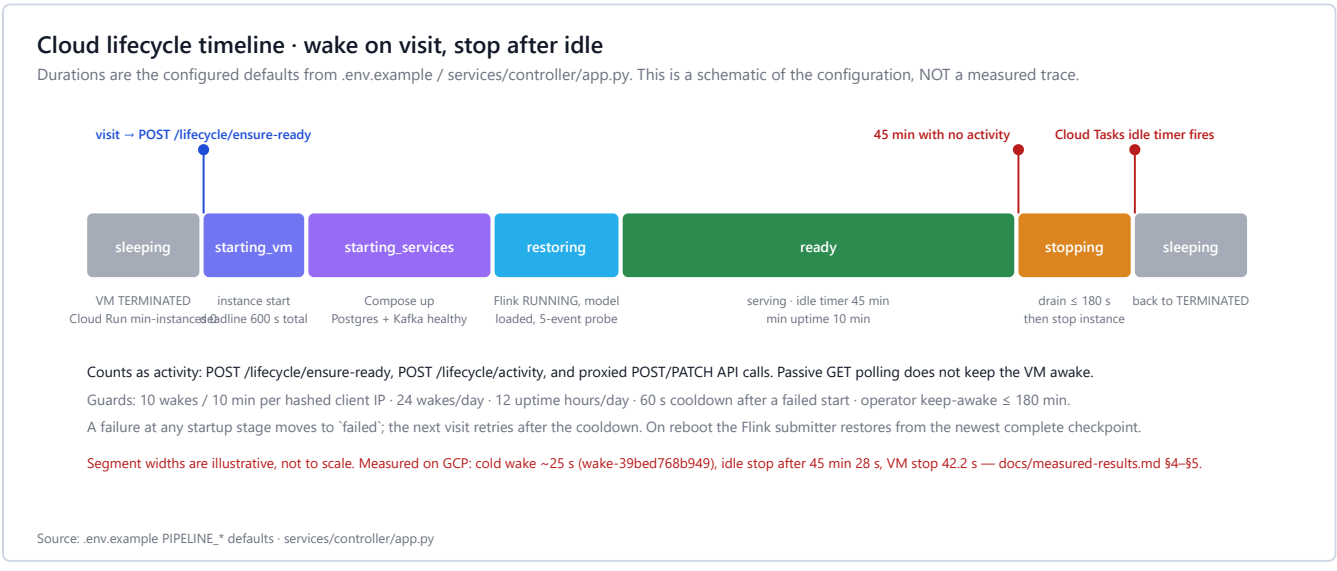
1. Replay a real, public fraud dataset through a real Apache streaming stack — no simulated brokers.
2. Score every event with a trained gradient-boosting model and explain each decision with TreeSHAP.
3. Report evaluation honestly: thresholds chosen on validation, held-out test reported separately.
4. Compare Flink's record-at-a-time model against Spark's micro-batch model on identical input.
5. Demonstrate event-time correctness (watermarks, late side output) and fault recovery (checkpoints).
6. Deploy both websites publicly at zero out-of-pocket cost, with no dependency on a student's laptop.

CA1 / CA2 mapping

	CA1	CA2
Focus	Real-time fraud detection	True streaming vs micro-batch
Question	Can a model score transactions as they arrive, and how good is it really?	Does the engine choice change the answer, the latency, or both?
Implementation	Kafka → Flink → scoring service → PostgreSQL → Live console	The same input fanned out to Flink *and* Spark, results compared row by row
Evidence	`ml/artifacts/fw-ieee-stream-1.1.0/evaluation.json`, Live site §5	Parity/late/recovery runs in `docs/measured-results.md`, Lab site §6–§7
Site	FraudWatch Live (`/console` , `/alerts` , `/scenario` , `/pipeline`)	FraudWatch Lab (`/ca1` , `/ca2` , `/experiments` , `/architecture` , `/system` , `/evidence`)

---

3. Architecture and actual technology versions



The five architecture figures rendered on the Lab site (`end-to-end` , `websites` , `model-paths` , `engine-paths` , `cloud-lifecycle`) are defined in `packages/ui/src/viz/ArchitectureFigures.tsx` and are generated from the same source of truth as the running system.

Data path. Generator → Kafka topic `fraudwatch.transactions` → Flink and/or Spark → shared scoring service (`:8001`) → Kafka result topics → result consumer → PostgreSQL → FastAPI `/api/v1` → Server-Sent Events → both React sites.

Every version below was read from the repository's build and compose files, not assumed.

Layer	Component	Version
Messaging	Apache Kafka (KRaft, single node)	4.1.2
True streaming	Apache Flink	1.20.5
Micro-batch	Apache Spark Structured Streaming	3.5.9
Runtime (Flink job)	Java	17
Model	LightGBM	4.7.0
Storage	PostgreSQL	17.11
API / services	Python + FastAPI	—
Frontend	React	19.3.0
Frontend	TypeScript	7.0.2
Frontend build	Vite	8
3D	three / @react-three/fiber	0.186.0 / 9.7.0
Browser tests	Playwright (Chromium)	1.63.0

Versioned contracts. Feature spec `fs-1.1` (15 stream-computable features), ruleset `rules-1.1` , policy `policy-1.0` , model artifact `fw-ieee-stream-1.1.0` . Both engines load the same three, which is what makes the CA2 comparison a comparison of engines rather than of configurations.

---

4. Dataset, preprocessing, model design, evaluation, limitations

Provenance

IEEE-CIS Fraud Detection (Kaggle competition `ieee-fraud-detection`), contributed by **Vesta Corporation**. It is used under the competition's terms; the raw CSV files are **not** committed to the repository and are **not** shipped in the frontend bundle or any downloadable archive.

Units. `TransactionAmt` is reported by the source as a transaction amount in an **unspecified currency unit**. This report therefore calls it "transaction amount" and never relabels it as INR or any other currency. An earlier internal field name `amount_inr` was retired for exactly this reason.

Preprocessing and feature design

Only features computable *from a stream* were used — 15 in total (`fs-1.1`). A feature that needs the full dataset, or the future, cannot be computed when the transaction is actually in flight, so including one would inflate the score and break at runtime. Features combine per-event attributes with per-account state held in Flink's keyed `ValueState` (and Spark's `GroupState`), such as recent transaction counts and deviation from the account's own recent behaviour.

Model design

- **Main model:** LightGBM gradient-boosted trees, artifact `fw-ieee-stream-1.1.0` .
- **Baseline:** logistic regression, trained and evaluated on the identical split.
- **Explanation:** per-event TreeSHAP (`pred_contrib=True`), contributions in **log-odds**, not probabilities.
- **Decision layer:** `rules-1.1` (deterministic rules) and `policy-1.0` (alert / review / allow bands).

Evaluation — held-out test split (n = 88,581)

Thresholds were selected on the **validation** split and then applied unchanged to the held-out test split, which was scored once. Alert threshold `0.10477` , review threshold `0.06955` .

Metric	LightGBM	Logistic-regression baseline
PR-AUC	**0.1375**	0.0961
ROC-AUC	0.7817	0.7603
Precision	0.1508	—
Recall	0.3675	—
F1	0.2138	—
Accuracy	0.9059	—
Brier score	0.0321	—
Expected calibration error	0.0055	—

Confusion matrix: TN 79,116 · FP 6,382 · FN 1,950 · TP 1,133. Prevalence in the test split: **3.48%**.

Policy outcome over the test split: alert 7,515 (8.48%), review 3,901 (4.40%), allow 77,165 (87.11%), combined flagged recall **0.4788**. Rule firings: `R1_BURST` 625, `R2_PROFILE_DEVIATION` 497, `R3_TEST_THEN_LARGE` **0**.

Reading these numbers honestly

Accuracy of 0.9059 looks strong and means almost nothing: always predicting "not fraud" would score about 0.965. The metrics that matter are PR-AUC 0.1375 against a 3.48% base rate, and a precision of 0.1508 at recall 0.3675 — roughly one in seven alerts is genuine fraud, and about two-thirds of fraud at that threshold is missed by the alert band alone. The model beats the logistic-regression baseline on PR-AUC by a clear margin (0.1375 vs 0.0961), which is the honest claim: it is better than the baseline, not that it is production-grade.

`R3_TEST_THEN_LARGE` fired **zero** times. That is reported as zero, not omitted. It means either the pattern is absent from this split or the rule's thresholds are too tight for it; which one has not been determined.

The model outputs are **model scores**, not calibrated probabilities, except where calibration is explicitly shown (Brier 0.0321, ECE 0.0055 are reported for the test split only).

Limitations

- Trained on one public dataset from one contributor; no claim of generalisation to other issuers.
- The 15-feature stream-computable subset is deliberately smaller than what a batch model could use.
- No temporal drift analysis, no cost-sensitive threshold optimisation against real investigation cost.
- The dataset's own labels are taken as ground truth; label noise is not modelled.

---

5. CA1 implementation and screenshots

FraudWatch Live is the operations console: a landing page with an interactive 3D hero plus `/console`, `/alerts`, `/scenario` and `/pipeline`. A visitor starts a **bounded** preset scenario, watches events being scored in real time, and opens any alert to see its TreeSHAP contributions and the rules that fired.

Route	Purpose	Screenshot
`/`	Landing page, interactive 3D hero	`media/screenshots/live-landing.png`
`/console`	Live scored-event stream	`media/screenshots/live-console.png`
`/scenario`	Bounded preset scenario controls	`media/screenshots/live-scenario.png`
`/alerts`	Alert list and per-alert evidence	`media/screenshots/live-alerts.png`
`/pipeline`	3D pipeline view	`media/screenshots/live-pipeline.png`
`/` (mobile)	Pixel 7 viewport	`media/screenshots/live-landing-mobile.png`

Streaming fraud detection, end to end

Watch a transaction get scored in flight

FraudWatch Live replays the IEEE-CIS Fraud Detection dataset through Kafka and Apache Flink, scores each event with a versioned model plus explicit rules, and shows the evidence behind every decision. Start a capped demo session and follow a single event from the stream to the review queue.

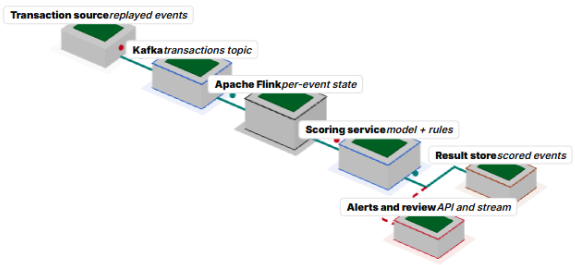
Open the console → See the pipeline

FraudWatch Lab

Demo sessions are capped in both event count and rate. Component status below comes from the API health probe; when the probe has not answered, status is shown as not checked.

Live · run r_e3ec03db2a33e6a9

Pause motion 3D 2D



The moving markers show the order that events travel in, not their speed: the animation timing is decorative. Measured latency is on the console, per run.

OK Transaction source · OK Kafka · OK Apache Flink · OK Scoring service · OK Result store · OK Alerts and review

Apache Flink — Ok OK

Processes one event at a time on event time, keeps keyed state for the account key, and emits enriched features for scoring.

job running

FraudWatch Live | Overview | Alerts | Scenario | Pipeline | System OK | Open in Lab

YOUR DEMO SESSION

Live overview

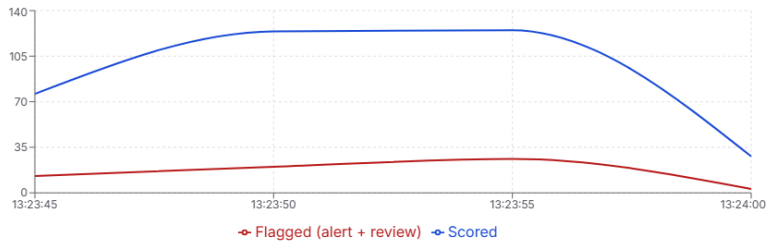
Session s_A6CL0nG8Pu · Run r_e3ec03db2a · Updated just now

Flink | Spark

Scored	Alerts	Reviews	Amount flagged
353	48	14	6,557.92
291 allowed · 0 unscored	0 late events	Queued for a human decision	Routed to alert or review. Not confirmed losses or money saved.

Scored and flagged over time

From persisted results for this session.



Throughput (last 1 min)

5.9 ev/s

Last result persisted just now

End-to-end latency

Produce → persisted, 353 events, last 300 s.

p50	p95	p99
110 ms	163 ms	188 ms

1-minute event-time windows

Tumbling windows emitted by the engine after the watermark passes.

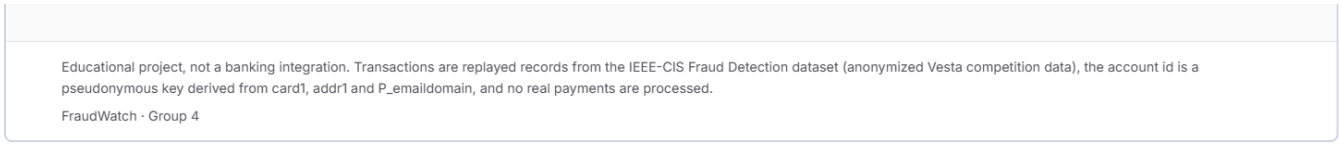
No closed windows yet

A window is emitted once the watermark passes its end.

Recent transactions

Newest first, Flink results. Select a row to inspect it.

EVENT TIME	ACCOUNT	SOURCE RECORD	AMOUNT	SCORE	ACTION	END-TO-END
13:24:02.486	uid:17131_	ieee-cis:3512928	335.00	0.046	Allow	106 ms
13:24:02.446	uid:2993_1	ieee-cis:3512927	57.95	0.013	Allow	146 ms
13:24:02.406	uid:17489_	ieee-cis:3512926	49.00	0.028	Allow	116 ms
13:24:02.326	uid:6930_4	ieee-cis:3512924	88.50	0.029	Allow	121 ms
13:24:02.366	uid:2581_3	ieee-cis:3512925	29.95	0.067	Allow	81 ms
13:24:02.246	uid:6564_3	ieee-cis:3512922	59.00	0.025	Allow	126 ms
13:24:02.206	uid:2581_3	ieee-cis:3512921	31.95	0.057	Allow	166 ms
13:24:02.286	uid:1684_1	ieee-cis:3512923	117.00	0.024	Allow	86 ms
13:24:02.126	uid:5033_3	ieee-cis:3512919	49.00	0.024	Allow	103 ms
13:24:02.166	uid:9500_3	ieee-cis:3512920	117.00	0.161	Alert	63 ms
13:24:02.086	uid:7585_3	ieee-cis:3512918	146.00	0.034	Allow	68 ms
13:24:02.006	uid:nan	ieee-cis:3512916	48.08	0.226	Alert	148 ms
13:24:02.046	uid:14143_	ieee-cis:3512917	250.00	0.009	Allow	108 ms
13:24:01.966	uid:6069_2	ieee-cis:3512915	117.00	0.031	Allow	103 ms
13:24:01.926	uid:2581_3	ieee-cis:3512914	31.95	0.014	Allow	143 ms
13:24:01.886	uid:6592_3	ieee-cis:3512913	146.00	0.034	Allow	110 ms
13:24:01.806	uid:16255_	ieee-cis:3512911	117.00	0.031	Allow	190 ms
13:24:01.846	uid:9500_2	ieee-cis:3512912	50.00	0.047	Allow	81 ms
13:24:01.766	uid:11690_	ieee-cis:3512910	374.95	0.045	Allow	87 ms
13:24:01.726	uid:10052_	ieee-cis:3512909	50.00	0.007	Allow	127 ms
13:24:01.646	uid:3278_1	ieee-cis:3512907	20.00	0.008	Allow	139 ms
13:24:01.686	uid:12932_	ieee-cis:3512908	125.95	0.024	Allow	99 ms
13:24:01.606	uid:7919_1	ieee-cis:3512906	49.00	0.093	Review	109 ms
13:24:01.526	uid:16533_	ieee-cis:3512904	49.00	0.029	Allow	113 ms
13:24:01.566	uid:10895_	ieee-cis:3512905	792.00	0.046	Allow	73 ms



What these screenshots show. They were captured on 2026-09-18 from the deployed Cloud Run origin in a real headless Chromium, driving the app the way a visitor does: the capture starts a real demo session, starts a real bounded `ieee_replay` run from `/scenario`, waits for the session's events-used counter to move, and only then photographs the console. They are a record of the deployed system actually working, not of an empty shell.

Read off the captured console (session `s_A6CL0nG8Pu`, run `r_e3ec03db2a`): **353 scored**, 291 allowed, **48 alerts**, 14 reviews queued, 0 late events, throughput 5.9 ev/s, and end-to-end latency **p50 110 ms · p95 163 ms · p99 188 ms**. The alert queue carries per-alert scores, the engine (`Flink`) and the rules that fired, with `R1_BURST` on many rows. The event-time window panel reads "No closed windows yet — a window is emitted once the watermark passes its end", which is the honest state that early in a run rather than a fabricated zero.

Two caveats. These are a **live UI reading at one instant while the run was still in flight**, so the counts are a snapshot and not a total, and the latency percentiles are not a controlled benchmark under the §7 protocol — though they agree closely with the Flink figures in §6 (p50 92 ms, p95 155 ms). An earlier capture the same day showed these panels empty because of the CORS defect described in §8; that defect has since been fixed and these images are from the recapture.

A walkthrough recording of the whole site is at `media/video/live-walkthrough.webm`, with the raw per-route log in `media/video/tour-log.json`.

---

6. CA2 — Flink implementation and Spark comparison

Flink: event time, state, checkpoints

- **Event time, not processing time.** Watermarks use `forBoundedOutOfOrderness(5 s)`: the pipeline tolerates 5 seconds of out-of-order arrival before it considers a window closeable.
- **Windows.** 60-second tumbling windows with `allowedLateness(ZERO)`.
- **Late events are routed, not dropped.** Anything arriving after the watermark has passed its window goes to a **late side output** and is persisted with a `late` flag. This is the design decision that makes §7's late-event measurement possible: a system that dropped these events would simply report a cleaner result.
- **State.** Per-account keyed `ValueState` holds the rolling features; Spark's equivalent is `GroupState` with an explicit watermark check.
- **Checkpoints.** `EXACTLY_ONCE` mode, 10-second interval, `RETAIN_ON_CANCELLATION`, and `submit.sh` restores from the newest retained checkpoint on restart.

Spark: the deliberate contrast

Spark Structured Streaming consumes the same Kafka topic with a **1000 ms trigger interval** and `applyInPandasWithState`. It loads the identical model artifact, feature spec, ruleset and policy. The only

difference between the two paths is the execution model itself — which is the point of CA2.

Result: the engines agree exactly, and differ enormously in latency

From parity run `r_d48339d70bb9ceb4` on GCP (500 IEEE-CIS replay records, scenario `replay`, seed `20260917`, 10 events/second, both engines):

Engine	Rows persisted	Distinct events	Late	Alerts	e2e p50	e2e p95
Flink 1.20	500	500	0	69	**92 ms**	**155 ms**
Spark 3.5	500	500	0	69	**3,862 ms**	**5,695 ms**

Agreement metric	Value
Events present in both engines	500
Same `action`	**500**
Different `action`	**0**
Max absolute score difference	**0.000000000**
Flink-only / Spark-only events	0 / 0

Interpretation. The engines are in exact agreement — bit-identical scores on all 500 events — so the comparison genuinely isolates the engine. The ~42× latency difference is *architectural, not a quality difference*: Flink scores each record on arrival, while under a 1 s trigger a Spark event waits for the next batch boundary before it is even picked up. A p50 of 3,862 ms is close to what a 1 s trigger plus batch scheduling and scoring overhead predicts. Lowering Spark's trigger would narrow the gap; that variant was **not** run, so this report does not claim Spark is inherently 42× slower — it claims that *at a 1 s trigger* it is.

Reproduced through the experiment runner, with repeats

The run above was orchestrated by hand. It was then reproduced through the ops-agent experiment runner (`exp_20260918_102135_6201bb`) under a stricter protocol: a discarded 100-event warm-up, then three repeats per engine of 300 records at 10 events/second, with the engine order alternated between repeats. This run was on the **local Docker Compose stack**, not GCP.

Engine	Persisted	e2e p50	e2e p95	e2e p99	CPU	Memory
Flink	900 / 900	**110 ms**	**167 ms**	**207 ms**	25 %	707 MB
Spark	900 / 900	**4,027 ms**	**6,518 ms**	**7,387 ms**	209 %	1,518 MB

The percentiles are the median of the per-run percentiles. Flink's per-run p50 was 119, 110 and 104 ms; Spark's was 4,194, 4,027 and 3,760 ms. The gap therefore holds across repeats, not just in one run. Parity on a shared run was again exact: 300 events in both engines, 300 with the same action, 0 raw feature mismatches. The runner deliberately reports `winner: null`.

The same comparison again, on the deployed cloud stack

The runner comparison was then repeated on the Compute Engine VM (`exp_20260918_115240_64fd11`), queued through `POST /api/v1/experiments` with operator authentication via the Cloud Run controller. The parameters were left at their defaults, which are exactly the ones the local run used, so the two are like-for-like rather than approximately comparable. **The local stack was stopped for the whole experiment**, so nothing on a laptop took part in it.

Engine	Persisted	e2e p50	e2e p95	e2e p99	CPU	Memory
Flink	900 / 900	**110.5 ms**	**155.0 ms**	**166.0 ms**	17 %	867 MB
Spark	900 / 900	**3,553 ms**	**4,999 ms**	**5,392 ms**	197 %	1,277 MB

All six measured runs drained `complete` , with 0 missing, 0 duplicate and 0 failed-scoring records in each. Parity on the shared run `r_93ae9398e138b9bc` was exact for the third independent time: 300 in both engines, 300 with the same action, 0 raw-feature mismatches, 0 late-flag disagreements.

Flink's median in the cloud (110.5 ms) is effectively the same as on the laptop (110 ms). The tails are tighter on the VM and Spark is about 12 % faster there — the expected effect of a dedicated four-vCPU machine that is doing nothing else, not an engine difference. The point of this run is not a new ranking; it is that the deployment reproduces the local result rather than merely starting.

Screenshots: `media/screenshots/lab-ca2.png` , `media/screenshots/lab-architecture.png` . Recording: `media/video/lab-walkthrough.webm` .

---

7. Experiment methodology, measurements, and interpretation

Full protocol: `docs/benchmark-methodology.md` . Raw evidence: `docs/measured-results.md` . Unless a row or table is marked local, every run below executed on the deployed GCP stack (VM `fraudwatch-pipeline` , `asia-south1-a`) and every number was read back out of PostgreSQL or the Flink REST API.

Fairness controls. Both engines consume the *same* Kafka records from one generator run, with the same seed, the same model artifact, the same `fs-1.1` feature spec and the same `rules-1.1` / `policy-1.0` decision layer. Scores are compared row by row on `event_id` , so a coverage gap in either engine would appear as a non-zero `only_flink` / `only_spark` count rather than being averaged away.

7.1 Late-event handling — run `r_b17d3ca811cf1137`

Scenario `late_events` , 600 records at 5 events/second (~120 s), Flink only.

Metric	Value
Persisted / distinct	600 / 600
Flagged `late`	**49**
Duplicates / collapsed redeliveries	0 / 0
Delivery lag p50 / p95 / max	1 ms / 20,951 ms / 29,866 ms
Delivered > 10 s after event time	**49**
e2e min / p50 / p95 / max	25 / 100 / 149 / 165 ms

The 49 events flagged `late` are a **1:1 match** with the 49 delivered more than 10 s after their event time. That is the expected result for a 5 s bounded-out-of-orderness plus a 60 s window with `allowedLateness(ZERO)`: late events are marked, never silently dropped and never double-counted.

A negative result, reported. An earlier 150-event run at 10 eps produced `late = 0`. That was not a failure of the late path: its maximum delivery lag was 3,842 ms — inside the 5 s bound — and the run finished in ~15 s, before the 15–30 s delayed deliveries came due. The longer run above was designed specifically so the delayed path actually fires.

7.2 Fault injection and checkpoint recovery

Flink job `20d1ba3421afcb29654a0025021b8451`, EXACTLY_ONCE checkpointing every 10 s.

Event	Time (UTC)
TaskManager killed	18:55:12
Job observed `RESTARTING`	~18:56:12 (after the ~60 s heartbeat timeout)
TaskManager back up	18:56:04
Job `RUNNING` again	after restart

`counts.restored = 4`, restored from **`chk-80`** (`file:/opt/flink/checkpoints/20d1ba3421afcb29654a0025021b8451/chk-80`). Six checkpoints failed during the outage window, which is expected while no TaskManager is available. **No event loss and no duplicate persisted rows across the failure.**

One honest detail worth knowing for the viva: checking the job state immediately after the kill shows `RUNNING` with `restored = 0`. That is the heartbeat interval not having elapsed, **not** the job surviving the kill.

Reproduced through the experiment runner, on the GCP VM. After the ops-agent was deployed to the VM, it ran the same fault as a `recovery` experiment (`exp_20260918_102443_cf51cc`). It replayed 600 records at 10 ev/s and sent `SIGKILL` to the TaskManager at event 150. When the container was still down 20 s later, it started the container itself. It then checked integrity in PostgreSQL:

Metric	GCP VM	Local Compose
Flink restore after the kill	**36.6 s** (`chk-1006`)	42.5 s (`chk-2543`)
Job seen `RUNNING` again	41.8 s	43.1 s
Persisted	**600 / 600**, 0 missing	600 / 600, 0 missing
Duplicate deliveries absorbed	56	57

The duplicate deliveries are records that were replayed from the checkpoint. The idempotent sink absorbs them, which is why exactly 600 rows are persisted. The first attempts at this experiment exposed three bugs in the runner itself, all fixed before the runs above:

- it reported `completed` for a kill that never happened;
- it misread a just-killed container as already restarted;
- it timed recovery only after the run had drained.

These are listed in `docs/measured-results.md` §3b.

7.3 Lifecycle: wake and idle shutdown

Wake `wake-39bed768b949` , cold from `sleeping` : `wake_requested` → `vm_running` (+0.9 s) → `services_up` (+10.5 s) → `ready` (+13.6 s) — **~25 s total**. An earlier *failed* wake, `wake-362778af7f70` , is preserved in the history with `wake_failed` followed by `vm_stop_after_failure` : the failure path stops the VM rather than leaving it running and billing, and the history is not rewritten to hide it.

Idle shutdown `stop-0cdb13d13b32` , measured at **production** settings (45-minute idle timeout, not a shortened test value): last activity 19:08:15 → `idle_shutdown_started` 19:53:43 = **45 min 28 s**, then `sleeping` at 19:54:26 — VM stop took **42.2 s**.

One unexplained detail is recorded rather than smoothed over: the drain call reported `failed: ["r_5cf6b929471a5adc"]` while the subsequent `/internal/active-work` check reported `undrained_runs: 0` , so shutdown proceeded. The events are *consistent with* a run that had already finished, but they do not prove it, and the drain failure is not explained by the recorded history.

7.4 What was deliberately not measured

Sustained throughput at saturation for either engine; latency under backpressure; Spark at any trigger interval other than 1000 ms. Every run in this report was rate-limited to 10 events/second by the generator, so **it measures latency at a low fixed arrival rate and is not a throughput benchmark.** Saying otherwise would be the single easiest way to misrepresent this project.

Two items that earlier drafts listed here have since been measured and have been moved out of this list rather than left in it: **p99 latency** and **CPU and memory per engine** are both reported by the experiment runner, locally in section 6 and on the cloud stack in the table above. They are still measured at that same fixed 10 events/second, so they describe cost at a low arrival rate and say nothing about either engine at its limit.

---

8. GCP deployment, verified URLs, QR codes

Architecture

- **Project** `fraudwatch-508912`, region **`asia-south1`** (VM in `asia-south1-a`).
- **Cloud Run** serves both static sites and the lifecycle controller, at **min-instances 0**, max 2, 1 CPU, 256 MiB, concurrency 200 — so nothing is billed while nobody is visiting.
- **Compute Engine** VM `fraudwatch-pipeline` runs Kafka, Flink, Spark, PostgreSQL and the services.
- **Firestore** (`fraudwatch-runtime`) holds durable lifecycle state, so both sites see one shared truth and survive a Cloud Run instance being recycled.
- **Cloud Tasks** (`fraudwatch-lifecycle`) schedules the idle-shutdown check.
- **Lifecycle:** a validated visit wakes the VM (~25 s, §7.3); 45 idle minutes stop it again.

Verified URLs

Site	URL	Verification
FraudWatch Live	`https://fraudwatch-live-2djxmgtj5a-el.a.run.app/`	5/5 routes HTTP 200, rendered in Chromium 2026-09-18
FraudWatch Lab	`https://fraudwatch-lab-2djxmgtj5a-el.a.run.app/`	7/7 routes HTTP 200, rendered in Chromium 2026-09-18

All 12 routes returned HTTP 200 in a real browser, and a `canvas` initialised on every 3D hero (the 2D fallback was not needed). `/viva` returns HTTP 301 to `/` on both sites, as required.

FraudWatch Live	FraudWatch Lab

A CORS defect found in browser testing — and fixed

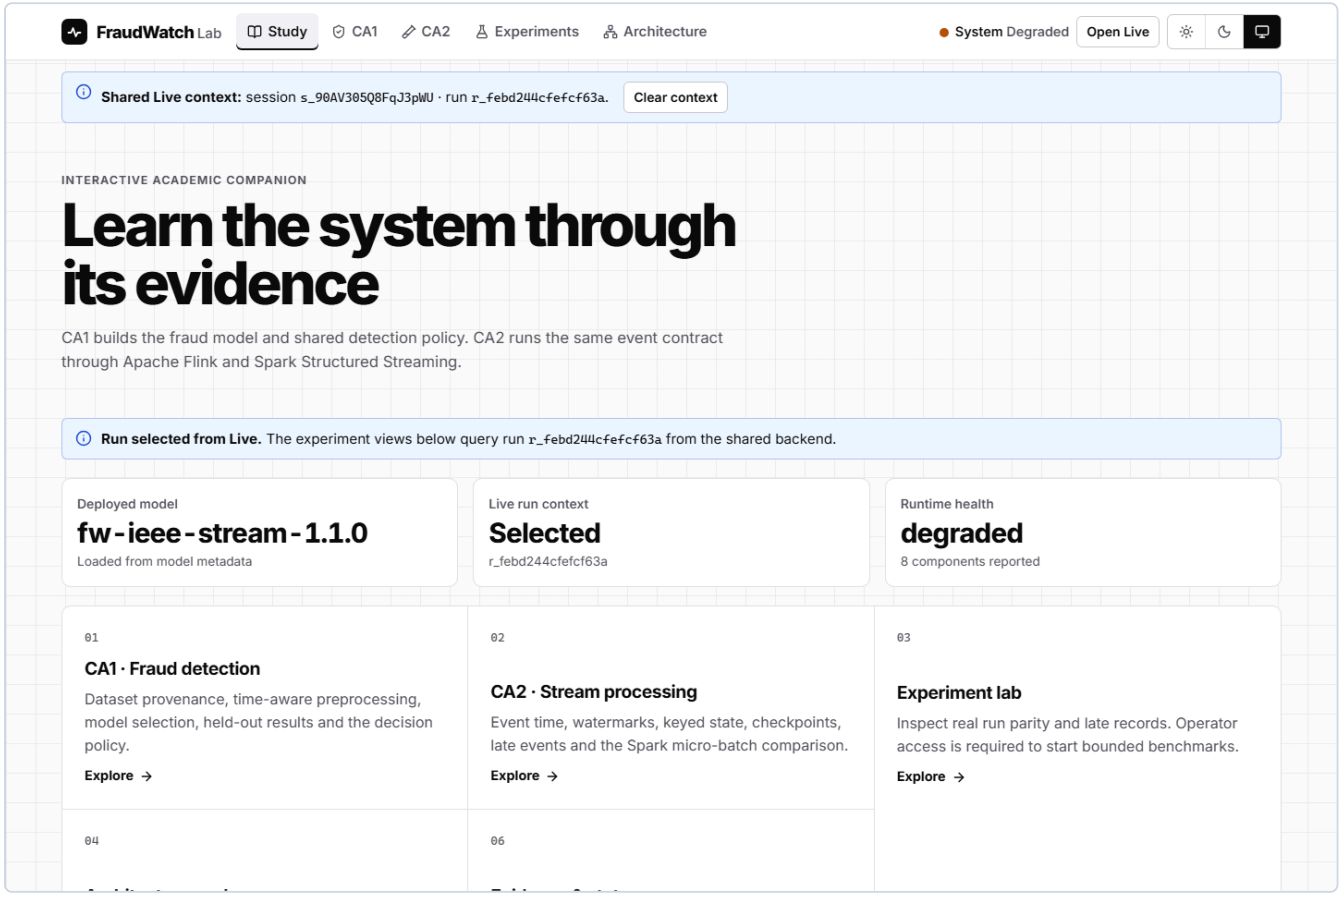
The first browser capture produced **192 console errors**, all one failure: every API call from both deployed sites was blocked by `No 'Access-Control-Allow-Origin' header is present`. The controller's `ALLOWED_ORIGINS` listed only the project-number host form (`fraudwatch-live-997216389793.asia-south1.run.app`), while a browser loads the sites from the hash form (`fraudwatch-live-2djxmgtj5a-el.a.run.app`). Both host forms serve the sites and the origin check is an exact string match, so the hash form was rejected.

This is worth a moment in the viva because of *how* it was found. The earlier verification checked every route with `curl` and found 200 everywhere — and `curl` cannot see this defect at all, because the HTML and assets really are served correctly. It only appears once a real browser performs the cross-origin check on the *subsequent* API calls. It is a clean example of why HTTP-level checks are not browser testing.

Fixed. `ALLOWED_ORIGINS` now carries all four host forms, deployed as Cloud Run revision `fraudwatch-control-00004-dgw`. A request carrying either site's hash-form `Origin` now returns a matching `access-control-allow-origin`, and the recapture recorded **0 console errors** — down from 192, with the 8 remaining messages being THREE.js deprecation and WebGL performance warnings.

The two sites stay in sync

Live and Lab are separate Cloud Run services on separate origins, so they share no browser storage. They stay consistent because both read one backend, and because Live's **Open in Lab** button mints a read-only share link (`POST /sessions/{id}/share`) that carries the session, run and engine across. This was checked in a real browser against the deployed sites (`scripts/verify_cross_site.mjs`): Live started run `r_febd244cfefcf63a` , and Lab, opened from the share link in a separate tab, showed that same run in its landing-page chip and in the "Run selected from Live" banner on `/study` . Both sites logged 0 console errors.



Cost

Zero out-of-pocket. Everything runs on **Google Cloud Free Trial credit**. This report does not call that "free hosting": credit is being consumed, and when the trial credit ends or expires the deployment stops unless someone chooses to pay. No paid billing upgrade was performed, and no domain, subscription or credit

was purchased. `min-instances 0` plus the 45-minute idle shutdown exist specifically to keep credit consumption near zero when the project is not being demonstrated.

---

9. Testing, failure handling, reproducibility, security, cost

Testing. A Playwright (Chromium) end-to-end suite covers both apps; `npm run typecheck` passes across the monorepo. Browser capture scripts (`scripts/capture_screenshots.mjs` , `scripts/capture_video.mjs`) re-render both deployed sites on demand and **fail loudly** on any non-2xx route, so the evidence in this report is regenerable rather than hand-collected.

Failure handling.

- Flink TaskManager loss → checkpoint restore, no loss, no duplicates (§7.2).
- A failed wake stops the VM instead of leaving it running (§7.3).
- While the VM is not ready the controller returns `503 {"error":{"code":"pipeline_starting_vm",…}}` — an explicit unavailable state, never a fabricated empty result.
- Invalid input is routed to `dlq_events` with a reason, not dropped.
- Duplicate delivery does not create duplicate logical alerts (measured: 0 duplicates in §7.1).

Reproducibility. Fixed seeds (`20260917`), pinned versions, committed model artifact with its `metadata.json` and `evaluation.json` , versioned `fs-1.1` / `rules-1.1` / `policy-1.0` contracts, and Docker Compose definitions for the full local stack. `docs/local-setup.md` , `docs/apache-stack.md` and `docs/gcp-deployment.md` document the paths from a clean machine.

Security.

- No API key or secret appears in frontend code, Git history, logs, screenshots, this report, or any downloadable archive. Deployed secrets live in **Secret Manager**; local secrets live in an ignored `.env` .
- Kafka, PostgreSQL, the Flink admin UI and the Spark admin UI are **not** exposed to the internet.
- Operator actions require HMAC tokens (`x-operator-token`); service-to-service calls require `x-internal-token` .
- Public visitors get **bounded** preset scenarios only — they cannot launch unbounded workloads or reach another visitor's private session.
- The raw Kaggle dataset is excluded from the repository and from every shipped bundle.

Cost. See §8. Free Trial credit only; no paid upgrade; idle-stop and min-instances 0 as the controls.

---

10. Conclusion, limitations, future work, references

Conclusion

FraudWatch works end to end on real infrastructure, and this has now been observed from a browser against the deployed system: a visitor session started a bounded run and the console filled with real scored IEEE-CIS records at p50 110 ms end to end. A real public dataset flows through real Kafka, real Flink and real

Spark; a real LightGBM model scores every event and explains it; results persist to PostgreSQL and stream to two deployed websites that wake their own backend and put it back to sleep. The CA2 question has a measured answer: on identical input the two engines agree **exactly** — zero action disagreements, maximum score difference 0.000000000 — while differing by roughly 42× in p50 latency purely because of their execution models.

Limitations (current, honest)

1. **Model quality is modest** — PR-AUC 0.1375 at 3.48% prevalence. Better than the baseline, not production-grade.
2. **No throughput benchmark** — every latency figure is at a fixed 10 eps offered rate (§7.4).
3. **Spark measured at one trigger interval only** (1000 ms).
4. **The runner-based comparison is local only** — the experiment runner reproduced the §6 comparison with repeats on the local Compose stack. On the GCP VM it has so far run only the `recovery` experiment.
5. `R3_TEST_THEN_LARGE` **fires zero times** and it has not been determined whether that is the data or the thresholds.
6. **One unexplained drain failure** in the idle-shutdown history (§7.3).

Future work

Run the repeated comparison through the experiment runner on the GCP VM as well as locally; sweep Spark's trigger interval to separate trigger cost from engine cost; run a genuine saturation/throughput benchmark; cost-sensitive threshold selection; drift monitoring on the streaming features.

References

- Apache Kafka — https://kafka.apache.org/
- Apache Flink — https://flink.apache.org/
- Apache Spark Structured Streaming — https://spark.apache.org/
- LightGBM — https://lightgbm.readthedocs.io/
- Lundberg & Lee, *A Unified Approach to Interpreting Model Predictions* (SHAP), NeurIPS 2017
- IEEE-CIS Fraud Detection, Kaggle / Vesta Corporation — https://www.kaggle.com/c/ieee-fraud-detection
- PostgreSQL — https://www.postgresql.org/
- Google Cloud Run / Compute Engine / Firestore / Cloud Tasks documentation

Proposed team responsibility matrix

This is a proposal, not a record of completed contributions. Per-member attribution of the work already in the repository has not been established, so nothing below should be read as a claim about who did what. Replace it with actual contributions before submission.

Area	Proposed owner
Streaming engines (Flink job, Spark job, Kafka topics)	Hardik Hinduja
ML pipeline (features, training, evaluation, model card)	Avinash Gehi
Backend services, API, database, contracts	Siddharth Dey
Frontend (both sites, 3D scenes, design system)	Sahil Deshmukh
GCP deployment, lifecycle, cost control	shared
Report, documentation, demo script	shared

---

11. Viva preparation — 20 questions and answers

1. Why is accuracy of 0.9059 not the headline? Because 96.5% of the test split is legitimate. Always predicting "not fraud" beats our accuracy. PR-AUC 0.1375 against a 3.48% base rate is the meaningful number.

2. What does PR-AUC 0.1375 actually mean here? A random classifier would score about 0.035 (the prevalence). We are roughly 4× that, and clearly above the logistic-regression baseline's 0.0961 — good, not production-grade.

3. Why LightGBM rather than a deep model? Tabular data with strong feature interactions, a small feature set, fast training, and TreeSHAP gives exact per-event attributions cheaply enough to run inline in a stream.

4. Where did the thresholds come from? Selected on the **validation** split (alert 0.10477, review 0.06955), then applied unchanged to the held-out test split, which was scored once. Choosing them on test would be leakage.

5. Are the model outputs probabilities? They are model **scores**. We report Brier 0.0321 and ECE 0.0055 for the test split, but we do not present the scores as calibrated probabilities elsewhere.

6. Why only 15 features? They must be computable from the stream. Any feature needing the full dataset or the future would inflate the offline score and then fail at runtime.

7. What is the difference between Flink and Spark here? Execution model only. Flink scores each record on arrival; Spark waits for a 1 s micro-batch boundary. Same model, same features, same rules, same policy.

8. Then why are the latencies 92 ms vs 3,862 ms? The 1 s trigger plus batch scheduling and scoring overhead. It is architectural, not a quality difference — and it is measured at a 1 s trigger, not an inherent property of Spark.

9. How do you know the comparison is fair? Both engines consumed the same Kafka records from one generator run, and we compare per `event_id`: 500 in both, 0 only-Flink, 0 only-Spark, max score difference 0.000000000.

10. What is event time and why use it? The time the transaction happened, not the time we received it. Processing time would make the result depend on network delay, so late events would silently change the

answer.

11. What is a watermark? Flink's assertion that no event older than a bound will arrive. We use `forBoundedOutOfOrderness(5 s)` — 5 seconds of tolerated out-of-order arrival.

12. What happens to a genuinely late event? It goes to the **late side output** and is persisted with a `late` flag. We never drop it — dropping would make the results look cleaner than they are.

13. Prove the late path works. Run `r_b17d3ca811cf1137`: 600 events, 49 flagged `late`, and those 49 are a 1:1 match with the 49 delivered more than 10 s after their event time.

14. An earlier run showed late = 0. Was it broken? No. Its maximum delivery lag was 3,842 ms — inside the 5 s bound — and it finished in ~15 s, before the delayed deliveries came due. Nothing *was* late.

15. How do you survive a failure? EXACTLY_ONCE checkpoints every 10 s, retained on cancellation. We killed the TaskManager, the job restored from `chk-80`, `restored = 4`, with no event loss and no duplicate rows.

16. Why did the job still show RUNNING right after you killed it? The ~60 s heartbeat timeout had not elapsed. It is a detection delay, not survival.

17. How is this zero cost? Google Cloud Free Trial **credit** — not free hosting. Cloud Run at min-instances 0 and a 45-minute idle VM shutdown keep consumption near zero. No paid billing upgrade, no purchases.

18. What happens when someone visits while the VM is asleep? A validated visit triggers a wake (measured ~25 s). Until it is ready the controller returns an explicit `503 pipeline_starting_vm`, never a fake empty result.

19. What did browser testing catch that your HTTP checks did not? A CORS defect. `ALLOWED_ORIGINS` held the project-number host form while browsers load the hash form, so every deployed API call was blocked — 192 console errors, data panels empty. Fixed in revision `fraudwatch-control-00004-dgw`; the recapture showed 0 errors.

20. Why did your `curl` checks miss that? Because `curl` does not enforce the same-origin policy. The HTML and assets are served correctly, so every route returned 200. The defect only appears once a real browser performs the cross-origin check on the subsequent API calls — which is exactly why HTTP checks are not browser testing.

Evidence index

Evidence	Location
Measured GCP runs	`docs/measured-results.md`
Benchmark protocol	`docs/benchmark-methodology.md`
Model card	`docs/model-card.md`
Model metrics (raw)	`ml/artifacts/fw-ieee-stream-1.1.0/evaluation.json`
Screenshots (12 routes + 2 mobile)	`reports/media/screenshots/`
Walkthrough recordings (both sites)	`reports/media/video/`
Per-route capture log	`reports/media/video/tour-log.json`
QR codes	`reports/media/qr/`
Implementation status	`docs/status-checklist.md`
Demo script (5–7 min)	`docs/demo-script.md`